

Induction + Orientation

Duration: 3 Days

Module 0

- Overview Of AI, Generative AI, And Agentic AI.
- Career Paths And Role Mapping For AI Professionals.
- LinkedIn Profile Optimization For AI & Tech Hiring.
- Using **Python Colab** For Hands-On Learning.
- Learning Workflow, Tools, And Program Onboarding.

Python for GenAI

Duration: 1.5 Months

Module 1

Topic 1: Core Python

- **Variables**
 - Store values in various data types
 - Dynamic assign / reassign variables
 - Python naming conventions for maintainable code
- **Data types**
 - Numeric types: **int**, **float**, **complex**.
 - Text handling using str and string methods.
 - Collections: **list**, **tuple**, **set**, **dict**.
 - Mutable vs immutable data behavior.
- **Loops**
 - Iteration using for and while loops.
 - Sequence traversal with **range()**, **enumerate()**, **zip()**.
 - Optimizing iterations for data processing tasks.

Topic 1: Core Python

- **Control Statements**

- Conditional logic using if, elif, else.
- **Comparison operators** (==, !=, <, >, <=, >=).
- Logical operators (and, or, not).
- Flow control using break, continue.
- Building **decision trees** for real-world logic handling.

- **Functions**

- Creating reusable logic blocks using def and return.
- Passing parameters using positional and keyword arguments.
- Anonymous function with **Lambda**

Project 1 Simple Calculator

1 hour

Create a console-based calculator that takes user input for two numbers and performs mathematical operations

Topic 2: Advanced Python

- **File Handling**

- Reading and writing files using open(), read(), write().
- Handling text, **CSV**, and **JSON** files for data pipelines.
- Working with file paths using os and pathlib.

- **Regex**

- Pattern matching using **re.search()**, **re.findall()**, **re.sub()**.
- Cleaning and extracting text data from logs, emails, and documents.
- Using regex tokens: \d, \w, +, *, ^, \$, groups.

- **Exceptional Handling**

- Custom exceptions handling using try, except, finally.
- Common exceptions.
- Preventing pipeline failures in production code.

Topic 2: Advanced Python

- **OOP's**

- Creating classes and objects using class and `__init__`.
- Instance variables vs class variables.
- Methods and object behavior modeling.
- Inheritance and method overriding.

Project 2 Banking System Simulation

BFSI

Simulate a simple banking system where users can create accounts, deposit, withdraw, and check balances.

Python

Data Structure

Function

Project 3 Movie Night Recommender

Entertainment

Recommend movies based on genre preferences using simple rule-based logic and conditional checks.

Data Types

Filtering Logic

User Input Handling

Project 4 Smart Fitness Habit Tracker

Healthcare

Track daily exercises, calories burned, and generate habit summaries using basic loops and functions.

Python

Functions

Conditional logic

Topic 3: Essential Libraries for GenAI

- **Numpy**
 - Creating arrays **using array, arange, linspace**.
 - **Array operations**: indexing, slicing, reshaping.
 - Vectorized computations for performance.
 - **Mathematical functions**: mean, sum, std, dot.
 - Broadcasting and handling multi-dimensional data.
- **Pandas**
 - Working with Series and DataFrame.
 - Reading data from **CSV, Excel, JSON**.
 - Data cleaning: missing values, duplicates, type casting.
 - Filtering, sorting, and conditional selections.
 - **Grouping and aggregation** using **groupby**.
 - Feature preparation for ML and GenAI pipelines.
- **Matplotlib**
 - Creating line, bar, scatter, histogram plots etc.
 - Customizing labels, titles, legends, and axes.
 - **Visualizing** trends and distributions in datasets.
- **Seaborn**
 - **Statistical plots**: boxplot, violinplot, countplot.
 - Relationship analysis using **pairplot** and **heatmap**.
 - Visualizing correlations and feature importance.
 - Styling plots for reports and dashboards.

Project 5 Retail Sales Data EDA

Retail

Analyze a retail sales dataset: clean data, plot trends (time series, bar charts), and derive insights.

Python

Pandas

Geopandas

Matplotlib

Topic 1: ML Fundamentals

- **Basic of ML**
 - **Supervised** vs **Unsupervised** learning paradigms.
 - Features, labels, training, validation, and test sets.
 - Bias-variance trade-off and model generalization.
 - **Overfitting** and **underfitting** with real examples.
 - End-to-end ML workflow used in industry projects.
- **Regression Analysis**
 - **Multiple Regression** with feature interactions.
 - Model evaluation using **RMSE**, **MAE**, **R² score**.

Project 6 Energy Consumption Forecasting

Energy

Predict household energy consumption using linear regression by analyzing temperature, appliances, and daily usage patterns.

Feature engineering

RMSE/R² evaluation

- **Classification Analysis**
 - **Logistic Regression** for binary classification.
 - Handling class imbalance using sampling techniques.
 - **Evaluation metrics: Precision, Recall, F1-Score, ROC-AUC.**

Project 7 Loan Approval Prediction

Finance

Build a classification model (e.g. logistic regression, random forest) to predict loan approvals from applicant data.

scikit-learn

Pandas

Topic 2: Deep Learning Fundamentals

- **Perceptron**
- **Backpropagation and Optimization**
- **Attention Mechanism**

Project 8 Smart Document Relevance Classifier

HR

Compute attention weights on simple sentences to understand token importance during context interpretation.

Numpy

Perceptron

Topic 3: NLP Fundamentals

- **Text Pre-Processing for Gen AI**
- **Text Embedding Techniques**
 - Bag-of-Words representation.
 - TF-IDF vectorization.
 - **Word2Vec: CBOW and Skip-Gram.**
- **Named Entity Recognition (NER)**
- **Topic Modeling**
 - Latent Dirichlet Allocation (LDA).
 - Latent Semantic Techniques (LSA/LST).
- **Part-of-Speech Tagging (POS)**
- **NLP Evaluation Metrics**

Topic 3: NLP Fundamentals

- **Transformer - BERT Model**
 - **Transformer architecture:** Encoder-Decoder blocks.
 - Multi-head self-attention mechanism.
 - Positional encoding concepts.
 - BERT pre-training objectives: **MLM** and **NSP**.
 - Use cases of **BERT** in real NLP systems.

Project 9 Sentiment Analysis Using BERT

Marketing

Use pretrained BERT embeddings to convert sentences into vectors and classify sentiment using a simple classifier.

BERT

Manual tokenization

Dense classifier

Project 10 Legal Entity Extraction

Legal

Perform named entity recognition and classification in legal documents using spaCy, and visualize results

spaCy

Pandas

Matplotlib

Project 11 Customer Churn Prediction

Telecom

Predict churn using ensemble models + SHAP for explainability

scikit-learn

XGBoost

SHAP

Advanced Generative AI

Duration: 2 Months

Module 3

Topic 1: Fundamentals of GenAI

- Difference Between Predictive AI And Generative AI.
- Generative Modeling Concepts And Probability Distributions.
- Pre-training vs Fine-tuning Paradigms.
- Foundation Models And Their Role In Modern AI Systems.

Topic 2: Use cases

- **Text Generation** For Chatbots And Assistants.
- Document Summarization And Report Automation.
- **Code Generation** And Developer Productivity Tools.
- Content Generation For Marketing And Media.
- Knowledge Assistants For Enterprises.
- **Search Augmentation** And **Question Answering Systems**.

Topic 3: Prompt Engineering

- **Zero-shot, One-shot, And Few-shot** Prompting.
- Instruction-based Prompt Design.
- Chain-of-Thought Prompting Techniques.
- Prompt Templates And Reusability.
- **Prompt Evaluation** And **Iterative Optimization**.

Topic 4: LLM

- **GPT Family** (GPT-3.5, GPT-4, GPT-4o).
- **LLaMA, Claude** and **Open-Source LLM** Ecosystem.
- **Tokenization** and **Context Windows**.

Topic 4: LLM

- **Inference Parameters: Temperature, Top-p, Max Tokens.**
- Hallucination And Limitations Of LLMs.
- Cost And Latency Considerations In Production.

Project 12 AI-Powered Sales Email Writer

Sales

Use an LLM to generate personalized sales outreach emails using customer behavior and CRM attributes.

OpenAI

Prompt Templates

Langchain

Topic 5: Hugging Face

- **Transformers Library** For NLP And GenAI.
- Pre-trained Models Hub And Model Cards.
- **Tokenizers** And **Pipeline** APIs.
- Fine-tuning Models Using **Hugging Face Trainer**.

Project 13 Product Review Sentiment Classifier

Retail

Build an NLP model to classify customer product reviews as positive or negative (using word embeddings and an LSTM or transformer).

NLTK/spaCy

Hugging Face

NLP Libraries

Topic 6: RAG (Retrieval Augmented Generation)

- **Vector DB**
 - Embedding Generation Using OpenAI And Hugging Face Models.
 - Vector Stores: **FAISS, Pinecone, ChromaDB, Weaviate, Qdrant**
 - Similarity Search And Distance Metrics.
 - **Chunking Strategies** For Optimal Retrieval.
- **Multimodal RAG**
 - Text + Image Retrieval Pipelines.
 - Multimodal Embeddings (**CLIP, Vision Transformers**).
 - Use Cases In Documents And Media Search.
- **Graph RAG**
 - Knowledge Graph Construction From Documents.
 - Relationship-based Retrieval **Over Vector Search**.
- **RAG Evaluation Metrics**
 - Context Precision And Recall.
 - Answer Faithfulness And Groundedness.
 - **Latency And Cost Evaluation**.
 - Human-in-the-loop Evaluation Methods.

Project 14 Custom Chatbot with Retrieval

Tech Support

Build a domain-specific chatbot: fine-tune an LLM on company FAQs and implement retrieval-augmented generation (Chroma or Elasticsearch) to answer user queries.

OpenAI API

Langchain

ChromaDB

Project 15 HR Policy RAG Query Bot

HR

Create a RAG-based chatbot that answers employee HR queries using internal company policies and guidelines.

Vector Store

Langchain

RAGAs Evaluation

Topic 7: Langchain / Llama Index

- **Prompt Templates And Output Parsers.**
- Chains And Sequential Workflows.
- Agents And Tool Calling.
- Integration With Vector Databases And APIs.
- Embed and store documents in vector DBs (**Pinecone, FAISS, Chroma**, etc.)
- Retrieve relevant chunks based on a user query
- Format the retrieved chunks as context for the LLM

Topic 8: Fine tuning

- **Full Fine-Tuning** vs **Parameter-Efficient Fine-Tuning (PEFT).**
- **LoRA** And **QLoRA** Techniques.
- Instruction Fine-Tuning.
- Dataset Preparation And Labeling Remotely.
- **Overfitting Risks** And **Mitigation Strategies.**

Topic 9: Fundamental Generative Models: GAN And VAE

- **Generative Adversarial** Networks Architecture.
- Variational Autoencoders And Latent Space Learning.
- Use Cases In **Image And Data Generation.**

Project 16 Synthetic Defect Image Generation

Manufacturing

Generate synthetic defect images to augment limited manufacturing datasets for improved model training performance.

GAN

Pytorch

TensorFlow

Topic 10: Stable Diffusion

- **Diffusion** Model Fundamentals.
- **Text-to-Image** Generation Pipelines.
- Prompt Control And **Image Conditioning**.
- Fine-Tuning Diffusion Models (**DreamBooth**, **LoRA**).

Project 17 Stable Diffusion Image Generator

Marketing

Use a Stable Diffusion model to generate creative images from text prompts; experiment with fine-tuning or style transfer.

CUDA GPU

Diffusers library

Topic 11: Flux Framework

- Image And Creative **Generation Workflows**.
- High-quality **Diffusion-based Outputs**.

Topic 12: Azure OpenAI

- **Azure OpenAI** Service Architecture.
- Model Deployment And Endpoint Management.
- **Enterprise Security** And Compliance.
- Integrating Azure OpenAI With Applications.

Project 18 Enterprise Document Assistant

Enterprise IT

Build a secure document Q&A system using Azure OpenAI embeddings and enterprise-grade retrieval pipelines.

Azure OpenAI

Python SDK

LangChain

Agentic AI & Automation

Duration: 1.5 Months

Module 4

Topic 1: Understanding Agentic AI and Autonomous Agents

- Difference Between **LLM Applications** And Agent-Based Systems.
- Autonomous Decision-Making And Goal-Oriented Agents.
- **Single-Agent vs Multi-Agent** Architectures.

Topic 2: Key components of Agentic AI

- Agent Roles, Goals, And Task Decomposition.
- Planning, Reasoning, And Execution Loops.
- **Tool Calling** And External System Integration.
- **Memory Types:** Short-Term, Long-Term, And Episodic Memory.
- Feedback Loops And Self-Reflection Mechanisms.
- Failure Handling And **Recovery Strategies**.

Topic 3: AutoGen

- **Multi-Agent** Conversation Design.
- Role-Based Agent Collaboration.
- Task Delegation And Message Passing.
- **Use Cases** For Autonomous Workflows.

Project 19 AI Data Analyst Agent

Finance

Develop an autonomous agent (e.g. using LangChain or LangGraph) that connects to a database and answers data questions by generating SQL queries and charts.

LLM (GPT-4o/Gemini)

Mistral Agents

Topic 4: CrewAI

- Team-Based **Agent Orchestration**.
- Role Assignment And Hierarchical Agents.
- **Workflow Automation** Using Crews.

Project 20 **Quality Inspector Agent**

Manufacturing

Build agent team to analyze defect images, generate insights, and suggest corrective actions collaboratively.

Vision models

OpenCV tools

CrewAI

Topic 5: LangGraph

- **State-Based Agent** Workflows.
- Directed Graph Execution For Agents.
- Conditional Routing And **Branching Logic**.
- Multi-Step Planning And Execution Graphs.
- **Error Handling** And Retry Mechanisms.

Project 21 **RAG Product Query Bot**

E-Commerce

Build a RAG assistant that answers product comparisons and queries using catalog data and customer reviews.

FAISS / Chroma

OpeAI LLM

LangChain

Topic 6: MCP (Model Context Protocol) / ACP / A2A

- Model Context Protocol (**MCP**) Fundamentals.
- Agent Communication Protocol (**ACP**).
- **Agent-to-Agent** (A2A) Interaction Patterns.
- Context Sharing Across Distributed Agents.
- **Secure Context Exchange** Between Models.
- Enterprise Interoperability Standards For Agents.

Project 22 MCP Coding Assistant & Debugger

IT

Build MCP-powered coding assistant that reads code files, runs tests, debugs errors, and suggests fixes automatically.

MCP SDK

OpenAI LLM

AutoGen

Topic 7: No-Code/Low-Code AI Agent Basics

- Build AI agents using **drag-and-drop workflow** builders.
- **Automate** tasks, approvals, and decision flows without coding.
- Evaluate feasibility, cost, and business **ROI** of AI agents.
- Work confidently with tech teams to scale AI faster.

Project 23 No-Code AI Agent Workflow Builder

Digital Transformation

Create a no-code AI agent workflow that automates business tasks, approvals, and decision flows using visual builders and rule-based orchestration.

LLM API

Rule Engines

n8n / Make.com

Project 24 Interoperable Multi-Agent System using MCP / ACP / A2A

Enterprise AI

Design a distributed multi-agent system where agents communicate, share context, and collaborate securely using standardized protocols across tools and services.

MCP SDK

AutoGen

A2A Protocol

Topic 8: AWS Bedrock

- Foundation Models Available In Bedrock.
- Building Agents Using Bedrock APIs.
- Tool Invocation And Workflow Automation.
- **Security, IAM, And Enterprise Controls.**
- Integrating Bedrock With AWS Services.

Project 25 Autogen Vehicle Diagnostics Agent

Automotive

Create diagnostic agent that interprets sensor logs and suggests repair actions using reasoning and tool execution.

AutoGen

AWS Bedrock

Claude

Project 26 LangGraph Human-in-the-Loop Approval Workflow

Enterprise

Agentic workflow with interrupt nodes requiring human approval before critical actions (e.g. email send, DB write).

LangGraph

FastAPI

LLMOP's & Production Deployment

Module 5

Duration: 1 Months

Topic 1: LLM Lifecycle, Versioning & Deployment Basics

- End-to-end LLM lifecycle from development to production.
- Model and prompt versioning using evaluation traces.
- API-based LLM deployment using FastAPI and OpenAI / Azure OpenAI endpoints.
- LangSmith, Langfuse, Phoenix, Arize for observability

Project 27 Chatbot Deployment

Tech Support

Build an LLMOps workflow to deploy, version, and monitor a chatbot using automated evaluation and logging.

FastAPI

Docker

LangSmith

Topic 2: Containerization & Inference Serving

- Packaging GenAI applications for consistent deployment.
- Serving LLM-powered APIs using containers.
- Local and cloud-ready inference setups.
- vLLM — PagedAttention for high-throughput open-source model serving
- Text Generation Inference by HuggingFace for production model hosting
- Kubernetes basics for scaling LLM services in cloud environments

Project 28 Production-Grade LLM Application Deployment

IT

Deploy an enterprise-ready LLM application with API endpoints, versioning, monitoring, and controlled rollouts for real-world usage.

FastAPI

Docker

LangSmith

Topic 3: RAG & Agent Deployment Pipelines

- Deploying Retrieval-Augmented Generation pipelines reliably.
- Vector database operations for production workloads.
- Deploying agent workflows with tool calling and state handling.
- **LangGraph** Platform — cloud-hosted agent deployment with persistent state

Project 29 RAG Q&A Pipeline

Knowledge Management

Build a Retrieval-Augmented Generation (RAG) pipeline using a vector database and LLM for Q&A.

Pinecone

HuggingFace

LangChain

RAGAs

Topic 4: Monitoring, Cost Optimization & Reliability

- Tracking latency, token usage, and failures.
- Prompt evaluation and trace-based monitoring.
- Cost optimization, rollback, and safe production controls.
- OpenTelemetry integration for LLM tracing standardization
- Token budget management and prompt compression techniques

Project 30 Agentic AI Monitoring & Cost Optimization System

Deploy an agent-based system that monitors LLM usage, detects performance or cost anomalies, and recommends optimization actions in production environments.

AutoGen

LangGraph

LLM

Topic 1: Responsible AI Foundations & Risk Awareness

- Understanding bias, hallucination, and misuse risks in GenAI systems.
- Transparency, explainability, and accountability in AI decisions.
- Enterprise expectations for ethical AI adoption.
- AI Act overview — risk tiers, prohibited uses, compliance obligations, data principal rights, consent framework for AI apps

Topic 2: Safety Guardrails & Human-in-the-Loop Controls

- Input and output validation strategies.
- Policy enforcement, prompt restrictions, and fallback handling.
- Human approval loops for high-risk AI actions.
- NeMo Guardrails (colang rules) + Llama Guard 3 (open-source safety classifier)
- Prompt injection detection and defense patterns

Topic 3: Governance, Compliance & Audit Readiness

- Monitoring AI behavior in production.
- Audit logs, traceability, and decision documentation.
- Aligning GenAI systems with enterprise and regulatory standards.
- Model cards and AI system documentation standards

Project 31 LLM Toxicity Monitoring & Safety Guardrail Engine

Responsible AI

Create a safety pipeline to detect toxic content, enforce prompt rules, and generate safer LLM responses.

OpenAI

Perspective API

Rule-based guardrails

Program Electives

Choose any 2 electives

A. GenAI and Agentic AI for **BFSI**

B. GenAI for **Retail & ecommerce**

C. GenAI for **Manufacturing & Automotive**

D. GenAI for **Healthcare & Pharma**

E. GenAI for **Cloud & Devops**

F. GenAI for **Managers & Leaders**

G. GenAI for **Software Developers**

H. GenAI for **QA & Test Engineers**

Gen AI Specialization for BFSI Professionals

Duration: 1 Month

Elective A

BFSI Domain Fundamentals

- KYC → onboarding → transactions → credit → servicing → collections
- Financial services workflows: portfolio advisory → compliance → fraud → risk scoring
- Insurance workflows: claims → underwriting → policy servicing → grievance redressal
- Key BFSI metrics: default rate, fraud probability, risk score, NPA metrics, claim turnaround, compliance accuracy
- RBI AI/ML guidelines for financial institutions — guardrails for AI in lending

Project 1 GenAI Credit Underwriting Document Assistant

Credit Risk

Build LLM assistant that reviews financial documents and generates eligibility summaries for credit decisions.

Azure

LLM Summarization

Document Parsing

Project 2 GenAI Claims Processing Assistant

Insurance

Summarize claim documents, detect missing information, and generate verification checklists for adjusters.

LLM Review

Structuring Templates

PDF Processing

Project 3 Agentic KYC Document Verification Pipeline

Compliance

Multi-agent pipeline extracting, validating, and cross-checking KYC documents automatically.

LangGraph

Vision Models

Rule Engine

Project 4 Agentic Compliance Review Agent

Banking

Build agent that checks policy documents against regulatory guidelines and flags compliance deviations.

RAG

Agentic Workflows

RBI Rule Parsing

Project 5 Agentic Fraud Detection & Explanation Engine

Fraud Detection

Implement multi-agent workflow that detects anomalies, explains fraud likelihood, and drafts investigation reports.

CrewAI

LangGraph

Reasoning Tools

Project 6 GenAI Digital Banking Q&A Assistant

Digital Banking

Create LLM-powered chatbot that answers customer queries using banking FAQs and product details.

LLM Prompting

Retrieval

Conversational Pipelines

Gen AI Specialization for Retail & E-com

Duration: 1 Month

Elective B

Retail & E-com Domain Fundamentals

- Retail workflows:
- catalog → listing → pricing → promotion → order → delivery → returns
- E-commerce stack overview:
- search, recommendation, merchandising, demand planning, customer support
- Key KPIs: conversion, AOV, returns rate, search-to-cart %, fill rate, OTIF

Project 1 Visual Product Search & Description Generator

E-commerce

Build a multimodal app: upload product image → GPT-4o generates title, description, tags automatically.

GPT-4o Vision

LangChain

Streamlit

Project 2 GenAI Returns Insights & Recommendation Assistant

E-commerce

Analyze return reasons using GenAI and generate improvement suggestions for products, logistics, or packaging.

LLM Summaries

RAG

Sentiment Models

Project 3 **Agentic Customer Query Resolution Agent**

Customer Support

Build multi-agent workflow that resolves customer queries, checks order status, and performs automated follow-ups.

Autogen

API Tool Calling

LangGraph

Project 4 **Agentic Catalog Enrichment Workflow**

E-Commerce

Create multi-agent workflow that extracts attributes, enriches product data, and validates outputs automatically.

CrewAI

LangGraph

CSV Tools

Project 5 **Agentic Pricing Recommendation Agent**

Retail

Build an agent that analyzes demand signals and competitor pricing to generate dynamic price suggestions.

LangChain

Data Tool Calls

LLM Reasoning

Project 6 **GenAI Bundle Recommendation Engine**

Growth & Analytics

Generate product bundle ideas using purchase patterns, browsing behavior, and LLM reasoning for higher conversions.

LLM Reasoning

Association Rules

Behavioral Data

Gen AI Specialization for Manufacturing & SupplyChain

Elective C

Duration: 1 Month

Manufacturing & Supplychain Fundamentals

- Manufacturing workflows:
- production planning → machining → assembly → quality → packaging → maintenance
- Supply chain workflows:
- procurement → inventory → warehousing → shipping → logistics → last-mile
- Key metrics:
- OEE, MTBF, defect rate, cycle time, lead time, OTIF, fill rate

Project 1 Assembly Troubleshooting RAG Agent

Manufacturing

Build RAG system that retrieves fault, tool, and assembly instructions to help technicians resolve issues faster.

LangChain

Vector DB

LLM Retrieval

Project 2 GenAI Inventory Insights RAG Assistant

Supply Chain

Build a RAG assistant that answers inventory, stock-level, and replenishment queries using supply chain documents.

LangChain

Vector DB

Azure Search

Project 3 Agentic QA Inspection Workflow

Quality Control

Implement multi-agent system where inspectors analyze images, classify defects, and generate corrective actions.

CrewAI

Vision Models

LangGraph

Project 4 GenAI Logistics RCA Assistant

Logistics

Use GenAI to analyze delivery failures and generate likely root causes with corrective action suggestions.

LLM Reasoning

Document Embeddings

RCA Prompt

Project 5 Agentic Warehouse Task Planner

Warehousing

Build an autonomous agents system that assigns warehouse tasks using SKU priority and worker availability.

CrewAI

LangGraph

Scheduling Logic

Project 6 GenAI Predictive Maintenance Report Assistant

Manufacturing

Build GenAI tool that explains sensor anomalies and generates human-readable maintenance recommendations.

LLM Reasoning

Time-Series Extraction

Summarization

Gen AI Specialization for Healthcare Professionals

Duration: 1 Month

Elective D

Healthcare Domain Fundamentals

- Healthcare workflows: patient intake → diagnosis → treatment → documentation → billing
- Pharma workflows: R&D → clinical trials → manufacturing → safety → pharmacovigilance
- Key metrics: ADRs, trial duration, readmission rate, LOS, claim turnaround, safety compliance
- HIPAA/HL7 FHIR basics — privacy and interoperability standards for AI in healthcare

Project 1 **Agentic ICD/CPT Coding Assistant**

Healthcare

Build agent that extracts diagnosis/procedure codes from clinical notes using rule checks and LLM reasoning.

CrewAI

Rule-Based Validation

RAG optional

Project 2 **RAG-Based Drug Safety Signal Investigator**

Drug Safety

Build RAG system that identifies adverse drug reactions by retrieving evidence from case reports and literature.

LangChain

Vector DB

LLM Reasoning

Project 3 **GenAI Trial Protocol Summary Assistant**

Clinical Trials

Summarize complex clinical trial protocols into simplified inclusion criteria, procedures, and risk summaries.

JSON Output

LLM Summarization

Document Chunking

Project 4 **Agentic Differential Diagnosis Generator**

Diagnosis Support

Create agent that analyzes symptoms, retrieves guidelines, and suggests possible differential diagnoses safely.

CrewAI

RAG

Safety Filters

Project 5 **GenAI FDA/EMA Compliance Checker**

Compliance & QA

Build system that checks SOP content against regulatory standards and highlights missing compliance elements.

RAG + LLM

Validation Layers

Regulatory Documents

Project 6 **GenAI Patient Information Assistant**

Patient Support

Build chatbot that explains diagnoses, treatments, and care instructions in simple language for patients.

LLM Prompting

Safety Filters

Conversation Flow

Gen AI Specialization for Cloud & DevOps Professionals

DURATION: 1 MONTH

Elective E

Generative AI in DevOps

- Introduction to GenAI in DevOps, AI for IaC automation, AI-assisted CI/CD
- AI-Powered Monitoring & Incident Resolution
- Predictive Analytics & Resource Management
- GenAI for Container Orchestration — Kubernetes manifest generation, Dockerfile
- Disaster Recovery & Security Automation with AI
- Fine-tuning LLMs for DevOps tasks + OpenTelemetry-based LLM observability
- GitHub Copilot / AWS CodeWhisperer — AI coding assistants in DevOps pipelines

Project 1 CI/CD Pipeline Optimization Assistant

Cloud

Build GenAI tool that analyzes CI/CD logs, detects bottlenecks, and suggests optimized pipeline steps.

Jenkins

GitHub Actions

LLM Reasoning

Project 2 Agentic Incident Response & Runbook Execution Bot

Cloud Ops

Create multi-agent workflow that triages alerts, references runbooks, and recommends resolution actions.

CrewAI

Runbook RAG

Alert Manager Logs

Project 3 **AI Log Pattern Summarizer & Anomaly Detector**

Observability

Build GenAI pipeline that clusters log patterns, detects anomalies, and generates incident insights automatically.

ELK

Embeddings

LLM Summaries

Project 4 **GenAI Security Drift & Compliance Checker**

Cloud Security

Detect misconfigurations, compare cloud states with policies, and generate compliance violation summaries.

OPA

LLM Reasoning

Cloud Config APIs

Project 5 **AI-Driven Kubernetes Auto-Scaling Agent**

Cloud Infra

Agent monitoring cluster metrics and automatically generating + applying HPA/VPA manifests using LLM reasoning.

LangGraph

Kubernetes Python Client

GPT-4o

Prometheus

Project 6 **GenAI Cloud Infrastructure Cost Optimizer**

FinOps

Analyze cloud resource usage, identify waste, and generate cost reduction recommendations using LLM reasoning.

LLM Reasoning

AWS Cost APIs

LangChain

Gen AI Specialization for Managers & Leaders

Duration: 1 Month

Elective F

Generative AI For Managers

- Agile Foundations & Practices — Scrum, Kanban, Jira, Sprint Planning
- Tools & Practical Execution — Jira, OpenProject, Slack, Teams, Confluence
- GenAI for Digital Business Transformation, AI Governance
- Agent-based Decision Systems, AI ROI Evaluation
- Prompt-based business automation with ChatGPT Enterprise / Microsoft Copilot
- Building an AI business case: ROI estimation, risk-benefit framework for C-suite

Project 1 AI Sprint Planning & Jira Workflow Setup

Agile Delivery

Plan a full AI sprint, create Jira boards, define epics, and break user stories into actionable tasks.

Jira

User Story Creation

Sprint Planning

Project 2 AI Meeting Intelligence & Action Item Tracker

Productivity

Transcribe meetings, extract decisions and action items, assign owners, and sync outputs to project tools.

Whisper STT

LangChain

Jira Integration

Project 3 **GenAI KPI Insights & Forecast Dashboard for Managers**

Business Strategy

Build AI-driven KPI dashboard that summarizes trends, identifies risks, and suggests business actions for leaders.

LLM Summaries

KPI Templates

Forecast Models

Project 4 **GenAI Marketing Campaign Design & Optimization Assistant**

Growth Strategy

Use GenAI to design campaigns, segment audiences, and generate strategic recommendations for managers.

LLM Reasoning

Prompt Templates

Customer Segmentation

Project 5 **AI Risk Assessment & Mitigation Planning Framework**

Compliance

Identify operational and financial risks, generate AI-assisted mitigation plans, and produce risk heatmaps.

LLM Analysis

Risk Scoring Matrix

Compliance Checklists

Project 6 **GenAI Root Cause & Action Plan Generator**

Operations

Use GenAI to analyze incidents, identify root causes, and auto-generate structured action plans for managers.

RCA Tree

LLM RCA Prompts

Workflow Templates

Gen AI Specialization for Software Developers

Duration: 1 Month

Elective G

Topic 1: AI-Powered System Design

- Designing scalable GenAI & agentic applications — context routing, retrieval layers
- Architecture choices for multimodal, RAG, multi-agent workflows
- Latency, cost, safety, reliability considerations

Topic 2: Production Architecture & Infrastructure Patterns

- Event-driven AI patterns, Model gateways, Inference routers, Distributed retrieval
- Monitoring, observability, fallback systems, guardrail enforcement
- FastAPI + Docker + Kubernetes + vLLM serving patterns
- Gateway pattern — single API façade routing to OpenAI / Azure / local vLLM based on cost/latency SLAs

Topic 3: Backend Architecture for LLM Applications & RAG Microservices

- Backend Architecture, High-Performance API (Async, Batching, Streaming)
- Caching, Persistence & Storage for AI Microservices
- Distributed RAG & AI Microservice Patterns
- Semantic caching — cache LLM responses by embedding similarity, not exact key match

Project 1 AI Inference Gateway & Routing Layer

System Design

Create an inference router balancing requests across different LLMs based on cost and performance.

API Gateway

Routing Logic

Project 2 High-Performance LLM API Service

API Architecture

Build a low-latency API service delivering LLM responses using async execution and intelligent routing.

FastAPI

AsyncIO

Project 3 Intelligent Prompt Caching Microservice

Backend Systems

Implement a backend service that caches prompts and responses to reduce repeated computation and system load.

TTL Strategies

Hash Keys

Project 4 Multi-LLM API Gateway with Fallback

System Design

API gateway routing to GPT-4o → Azure fallback → vLLM local, based on latency budget and cost limits.

FastAPI

LiteLLM

Docker

vLLM